


Llama-3.2-1B, batch 1, BF16, P32/D128 on H100



[†] SGLang is baseline-subtracted server-request E2E; the other bars use matched in-process decode-step timing.